

# Comprehensive Deployment Architecture for a Legal AI Platform

## Architecture Overview

This platform is designed as a full-stack legal AI system that integrates document processing, advanced AI reasoning, and a user-friendly web interface, all backed by robust cloud infrastructure and CI/CD. It will ingest legal documents (e.g. PDFs of cases, statutes, etc.), convert them to machine-readable text, and index that data for retrieval. A large language model (LLM) – such as GPT-4 or a fine-tuned open model – serves as the reasoning engine, capable of answering legal questions and drafting documents with human-level proficiency (GPT-4 has demonstrated bar exam performance around the 90th percentile). Key components include:

**Document Ingestion Pipeline:** Converts PDFs and other files into text via OCR, cleans and chunks the text, then creates vector embeddings for semantic search. This forms a knowledge base of laws, cases, and regulations for Retrieval-Augmented Generation (RAG).

**AI Reasoning Engine:** A powerful LLM (backed by Azure OpenAI or custom deployment) augmented by tools (e.g. a “Legal MCP server” for structured reasoning). It can retrieve relevant law from the vector database, perform legal reasoning, and draft outputs (opinions, briefs, etc.) with citations to source material. Fine-tuning and Reinforcement Learning from Human Feedback (RLHF) are used to continually improve accuracy.

**Frontend Web Application:** Enables users to interact with the AI via a browser – including uploading PDFs, answering questionnaires, selecting jurisdictions or case types, and receiving answers or document drafts. The UI will be modern and intuitive, with dropdowns for case type/jurisdiction and rich text editors for drafting.

**Security & DevOps:** All secrets (API keys, endpoints) are stored in Azure Key Vault for security. The application and services run in Docker containers orchestrated in Azure (decision on Azure App Service vs Kubernetes detailed below). CI/CD pipelines (Azure DevOps YAML) and Infrastructure-as-Code (Terraform) automate builds, tests, and deployments. Monitoring and logging are in place for performance and cost tracking.

**Admin & Subscription Management:** An admin portal allows oversight of user activity, content, and model performance. Stripe integration provides subscription billing with tiered plans. The service is optimized to be affordable for pro se users, undercutting traditional legal research tools (Casetext plans started ~\$65–\$90/month, and Westlaw can cost hundreds per month). Users must consent to data collection for model improvement and accept disclaimers that AI outputs may contain errors.

Below, we detail each aspect of this architecture and deployment plan.

## Document Ingestion & Knowledge Base Construction

1. **PDF Processing and OCR:** The platform will accept PDF uploads (e.g. a user’s case documents or a large corpus of statutes). Many legal PDFs (especially scans of filings or old cases) require optical character recognition. We utilize Azure AI Document Intelligence (Form Recognizer) with its OCR “Read” model, which is optimized for extracting printed or handwritten

text from PDFs and images. This service can handle multi-page documents and outputs structured text. (Alternatively, an open-source stack using pdf2image + Tesseract OCR could be used, but Azure's managed service offers higher accuracy and scalability.) The OCR step yields the raw text of each document page.

2. Cleaning and Structuring Text: Extracted text is preprocessed and cleaned to improve quality. This involves standardizing encodings, removing artifacts (line numbers, stamps), handling special characters, and segmenting content logically. For example, a PDF of a court opinion might be split into sections (headnote, background, reasoning) and cleaned of page headers/footers. Consistent structure and formatting are crucial for downstream tasks. The ingestion pipeline also extracts metadata (e.g. document title, jurisdiction, date) to tag the content.

3. Chunking and Embedding: Each document is broken into manageable text chunks (e.g. by paragraph or section) for indexing. We determine an optimal chunk size (perhaps ~500 tokens) to balance context completeness with retrieval granularity. Using an embedding model (such as OpenAI's Ada or Azure's text embedding model), the pipeline generates a vector embedding for each chunk. These embeddings capture semantic meaning so the system can later find relevant snippets by similarity search. (Notably, embedding models are very cheap to run – e.g. only ~\$0.0001 per 1,000 tokens for OpenAI's Ada – so this is cost-effective at scale.)

4. Vector Database and Search Index: All chunk embeddings are stored in a vector database for fast similarity lookup. We can use Azure Cognitive Search with vector indexing (now supported in Azure Search), or a self-hosted vector DB like Postgres+pgvector or Pinecone. Each vector entry links to the original text chunk and its source info. We also index basic keywords (for a hybrid search approach), assigning a relevance score to each document chunk. This allows the system to retrieve the most relevant legal texts (cases, statutes, rules) when a user query or drafting task is performed. In essence, this is a legal knowledge base encompassing city ordinances, county regulations, state statutes, federal laws, case law precedents, procedural rules, etc., all vectorized for AI consumption. As new data is ingested, the index is updated incrementally.

5. Domain-Specific Processing: The ingestion pipeline can be configured by an admin to handle specialized content. For example, if uploading a body of case law, the system can auto-detect case citations and create a knowledge graph of citations. (In future, this enables a citator feature: flagging if a case has been overruled or criticized, similar to how Casetext added an "orange flag" if a case relies on an overruled opinion.) If loading court rules or forms, the pipeline can apply templates to parse them into structured data (for use in document drafting). All data passes through a quality check – e.g. extremely short or irrelevant chunks are filtered out, and duplicate texts are merged – to keep the knowledge base efficient.

6. JSONL Export for Fine-Tuning: Besides powering retrieval, the cleaned text can be used to fine-tune models. The pipeline can output training data in JSONL format (JSON Lines) as needed. For instance, it might generate Q&A pairs from a statute (question: "What is the penalty for X per the statute?", answer: the statute text) or extract case summaries. These JSONL files serve as input to model fine-tuning jobs (OpenAI fine-tuning or custom model training). Over time, as more documents are ingested (including user-provided documents and our curated legal corpora), we can periodically fine-tune the LLM on this growing dataset to improve its

domain expertise. We will leverage large public legal datasets (e.g. the Pile of Law, a 256GB open corpus of legal texts) to bootstrap the model's knowledge and ensure broad coverage of US legal domains (excluding international law as specified).

This robust ingestion system ensures the AI always has up-to-date, relevant source material at its fingertips. Clean, indexed data is the foundation for reliable retrieval-augmented answers.

## AI Model and Legal Reasoning Engine

1. Core LLM Choice: At the heart is a Large Language Model capable of sophisticated legal reasoning. We propose using GPT-4 via Azure OpenAI Service, as it has demonstrated exceptional performance on legal tasks (for example, GPT-4 passed the Uniform Bar Exam, including essays and multiple-choice, at approximately the top 10% of test-takers). GPT-4's ability to understand complex prompts and produce lengthy, coherent legal text is unparalleled. We can start with GPT-4 (8k or 32k context as needed) for generation. To reduce costs for simpler queries, we might also incorporate GPT-3.5-Turbo for less complex tasks or lower-tier users. Another option is fine-tuning an open-source LLM (like Llama 2 or a GPT-J variant) on legal data to use for certain tasks, but given the breadth of legal knowledge required, GPT-4's zero-shot capabilities are a huge advantage in the short term.

2. Retrieval-Augmented Generation (RAG): To maximize factual accuracy, the model will use Retrieval-Augmented Generation. This means when a user asks a question or requests a document, the system first uses the vector knowledge base to fetch the most relevant law/cases, and includes those snippets (with citations) in the prompt to GPT-4. The LLM then bases its answer or draft on that provided context, greatly reducing hallucinations and ensuring any specific legal references are grounded in real sources. This approach lets us combine GPT-4's reasoning with our database of statutes and case law. The result is an AI that can not only answer questions in plain English, but also cite the relevant code sections or precedents to back up its answers – a critical feature for legal users.

3. MCP Server and Agentic Reasoning: For complex analyses, we incorporate an MCP (Model-Controller-Protocol) server pattern, which is an emerging architecture for orchestrating LLM "agents." In practical terms, this is a microservice that can decompose complex queries into structured steps and tool calls. For example, TrueLaw (a legal AI company) describes an agentic architecture where a model router decides which models or tools to use, and an RLHF-tuned workflow guides the AI through multi-step tasks. Our implementation can leverage the open-source Cerebra Legal MCP Server, which provides "tools" for legal reasoning (like `legal_think` for step-by-step analysis, `legal_citation` for formatting citations, etc.). When the user poses a broad task (e.g. "Analyze my case and suggest the best legal strategy"), the MCP server could break this down – first have the LLM identify key issues and possible strategies, then retrieve specific precedents for each strategy, then evaluate outcomes. This agentic approach makes the AI more "strategic" and context-aware, not just a Q&A bot. It can also decide when a straightforward answer is possible versus when to engage in a lengthy research session (saving costs by not always querying the vector DB or external APIs unnecessarily).

4. Fine-Tuning and RLHF: We will continually refine the model for legal applications. Fine-tuning is done on domain-specific data (e.g. examples of well-written legal arguments, Q&A pairs from bar exam prep, etc.) to align the model's style and knowledge. More importantly, we implement

Reinforcement Learning with Human Feedback (RLHF) to improve the model's outputs over time. Lawyers or expert reviewers in the loop (including perhaps our own team during beta) will evaluate the AI's answers for correctness, relevance, and tone. For instance, if the AI draft misses a required legal element or cites a non-existent case, the reviewer corrects it. These comparisons (preferred vs. original output) train a reward model and policy so the AI learns to avoid those mistakes. Using RLHF has been key to aligning models like ChatGPT toward factual and helpful responses. In our context, RLHF (possibly facilitated via tools like Weights & Biases for experiment tracking) will help minimize hallucinations and bias, ensuring the AI's legal advice is reliable. We will track metrics like factual accuracy and citation precision in a feedback database. The system also includes an automated quality evaluator (similar to Google's approach): after the AI answers, an automated check or secondary model can score the answer on factuality and completeness, flagging low-confidence results for human review.

5. Guardrails and Policies: Because legal output is high-stakes, we enforce strict guardrails. The model will always cite its sources (statutes, case law, or uploaded documents) when providing any legal statement of fact. It will include disclaimers in its responses reminding users "I am an AI assistant and not a licensed attorney. Please verify all suggestions and citations." Additionally, we maintain a session log of all AI-user interactions for transparency. If a user asks for something potentially illegal or disallowed (e.g. advice to violate the law), the system will refuse or safe-complete in adherence to OpenAI content guidelines. All these measures are clearly disclosed to the user. The famous incident of lawyers sanctioned for submitting AI-generated fake citations (*Mata v. Avianca*) underscores why our platform must prioritize accuracy. Through RAG and RLHF, plus user-facing warnings, we mitigate the risk of "hallucinated" legal citations. Users are contractually required (via Terms of Service) to take responsibility for verifying the AI's output before relying on it, which protects both the user and our service.

6. Performance and Scaling: The AI engine will be deployed in a scalable manner. If using Azure OpenAI, we leverage the Azure OpenAI resource with sufficient throughput (and possibly enable caching of common queries to reduce latency/cost). If we host our own model in containers, we'll use Azure Kubernetes Service (AKS) with GPU-enabled nodes for the model server and autoscale the number of replicas based on load. For embeddings and smaller tasks, lightweight CPU containers can scale out as needed. By routing different tasks to different models (as TrueLaw's router does), we ensure efficient use of resources – e.g. quick questions may go to a smaller model vs. lengthy research to GPT-4. This model routing can substantially cut costs while maintaining quality.

In summary, the AI layer combines a knowledge-grounded LLM with advanced orchestration. It can answer user questions with cited sources, draft documents in proper legal format, and even suggest strategic options – essentially functioning as a knowledgeable legal assistant. With continuous learning from real case outcomes and user feedback, its capabilities will only improve over time.

## Frontend User Interface and Experience

The front end is a critical component, as it's the touchpoint for users (whether self-represented litigants or lawyers seeking quick research). We will develop a responsive web application, likely using a modern JavaScript framework (React or Angular) for a smooth user experience. This

frontend will be containerized (Docker) and served via Azure App Service or a static hosting + CDN, depending on final architecture.

1. User Onboarding and Dashboard: Upon login, users land on a dashboard showing their recent queries, drafts, and an overview of their subscription status. The UI will be clean and professional, with legal-themed design elements but an emphasis on simplicity. A new user will be guided through the process with tooltips (e.g. “Upload a document to get started” or “Ask a legal question in the chat”). We plan to implement role-based access – most users are “Clients” who can only access their own data, while internal staff or admins have higher privileges (discussed later).

2. Document Upload and Analysis: Users can upload PDFs or images (like a photo of a ticket) for analysis. The frontend provides an upload form tied to the ingestion pipeline described earlier. When a PDF is uploaded, the UI will show a progress bar as it’s being processed (converted via OCR, etc.). Once done, the system could display an overview of the document – e.g. if it’s a court opinion, show the case name, date, and a summary generated by the AI. The user can then query the AI about that document (e.g. “Explain this decision” or “Find key points relevant to my situation”). This tight integration of document intelligence means the platform can serve as an e-discovery or document review assistant: the user might dump a 100-page contract and ask, “Which clauses should I be concerned about?” and the AI will highlight sections, thanks to the earlier OCR and indexing.

3. Legal Query Chatbot: A core feature is a chat interface where users ask legal questions in natural language. This is similar to ChatGPT but tailored for law. The UI shows a conversation: user questions and AI answers with citations. We ensure the chat supports rich text (the AI’s answers will include formatted citations, bullet lists for steps, etc.). Code blocks might be used if showing statutory text or formatted regulations. Users can thumbs-up or flag an answer if something seems off, feeding our feedback loop. Importantly, users can toggle to “Show Sources” which highlights the excerpts from the legal texts that the AI used (providing transparency and enabling users to double-check the AI’s references). This chat could also be used for an interactive questionnaire: the AI might ask the user a series of questions to gather facts (just as a lawyer would) and then produce a recommendation or draft.

4. Case Type Selection and AI Strategy: To guide the AI’s reasoning, the UI includes a drop-down menu or wizard for case type/jurisdiction. For example, a user might select “Landlord/Tenant – Texas” as their matter type. This selection can constrain the AI to use the Texas Property Code and relevant Texas case law in its answers (we will tag our documents by jurisdiction to facilitate this). If the user isn’t sure, they can choose “Not sure – let AI decide,” in which case the AI (via the MCP server) will analyze the user’s description and classify the case itself. Additionally, we provide an option like “Explore Options”: the AI can present multiple legal routes (e.g. “Option 1: File a small claims case (pros/cons)... Option 2: Attempt mediation... Option 3: etc.”) and even indicate which is most cost-effective or likely to succeed. This addresses the request for the AI to pick the “most economically feasible” solution – essentially giving a strategy recommendation along with likely time/cost for each path. While AI won’t perfectly predict legal outcomes, it can at least lay out common approaches and factors (drawing on its knowledge of similar cases).

5. Document Drafting Tool: A standout feature is the AI Document Drafter. The user can choose

from a library of document templates (e.g. “Motion to Dismiss”, “Will”, “Contract”, etc.) or let the AI decide the appropriate document. The frontend presents a rich-text editor pre-populated with the AI’s draft. Crucially, before drafting, the user can specify jurisdiction and formatting requirements: for instance, select a federal court vs. a specific state court. The AI will format the document according to that court’s rules (margins, caption, citations in Bluebook format, etc.). This is feasible because the LLM can be prompted with style guidelines, and tools like Spellbook have shown AI can “automatically adhere to legal writing standards and formatting guidelines”. For example, if the user needs a brief in California state court, the AI will ensure the font, line spacing, and citation format meet California Rules – possibly by having stored prompt instructions or style sheets for each jurisdiction. The user can edit the draft as needed in the editor. We also include a “Validate” function that cross-checks citations in the draft against the knowledge base to ensure they are real and properly quoted (to avoid any *Mata v. Avianca* scenario). Finally, the user can download the document (as Word or PDF) or save it in the system.

6. User Guidance and UX: Since many target users are pro se (not legally trained), the UX will be friendly and educational. Jargon will be minimized or explained. For instance, if the AI’s answer uses a term like “summary judgment”, the UI could have a tooltip explaining it. We’ll incorporate a glossary and FAQs. The design will emphasize trust – e.g. showing source links, and maybe an “AI confidence” indicator or statement like “Sources provided. Please verify content.”. We also allow users to chat with a human lawyer (if we partner with any attorney network) for an additional fee, but the primary goal